

Math 7243 Machine Learning - Homework 3

For programming questions, you can only use numpy library.

Question 1. Softmax regression Recall the setup of logistic regression: We assume that the posterior probability is of the form

$$p(Y = 1|\vec{x}) = \frac{1}{1 + e^{-\beta^T \vec{x}}}$$

This assumes that $Y|X$ is a Bernoulli random variable. We now turn to the case where $Y|X$ is a multinomial random variable over K outcomes. This is called softmax regression, because the posterior probability is of the form

$$p(Y = k|\vec{x}) = \frac{e^{\beta_k^T \vec{x}}}{\sum_{j=1}^K e^{\beta_j^T \vec{x}}}$$

which is called the softmax function. Assume we have observed data $D = \{\vec{x}^{(i)}, y^{(i)}\}_{i=1}^N$. Our goal is to learn the weight β_1, \dots, β_K .

(1) Find the negative log likelihood of the data $l(\beta_1, \dots, \beta_K) = -\log L(\beta_1, \dots, \beta_K) = -\log P(Y|X)$

$$\begin{aligned} -\log \mathbb{P}(Y|X) &= -\log \prod_{i=1}^N \mathbb{P}(y_i|x_i) = -\log \prod_{i=1}^N \prod_{k=1}^K \left(\frac{e^{\beta_k^T x_i}}{\sum_{j=1}^K e^{\beta_j^T x_i}} \right)^{1\{y_i=k\}} \\ &= -\sum_{i=1}^N \sum_{k=1}^K 1\{y_i = k\} \left(\beta_k^T x_i - \log \left(\sum_{j=1}^K e^{\beta_j^T x_i} \right) \right) \\ &= -\sum_{i=1}^N \sum_{k=1}^K 1\{y_i = k\} \beta_k^T x_i + \sum_{i=1}^N \log \left(\sum_{j=1}^K e^{\beta_j^T x_i} \right) \end{aligned}$$

(2) We want to minimize the negative log likelihood. To combat overfitting, we put a regularizer on the objective function. Find the **gradient** w.r.t. β_k of the regularized objective

$$l(\beta_1, \dots, \beta_K) + \lambda \sum_{k=1}^K \|\beta_k\|^2$$

$$\nabla_{\beta_k} -\log \mathbb{P}(Y|X) = 2\lambda\beta_k - \sum_{i=1}^N 1\{y_i = k\}x_i + \sum_{i=1}^N \frac{e^{\beta_k^T x_i}}{\sum_{j=1}^K e^{\beta_j^T x_i}} x_i$$

Note that we can use the definition of $\mu_k(x_i)$ here to save a bunch of writing.

$$= 2\lambda\beta_k + \sum_{i=1}^N (\mu_k(x_i) - 1\{y_i = k\}) x_i$$

(3) State the gradient updates for both batch gradient descent and stochastic gradient descent.

Batch gradient descent:

$$\beta_k^{(t+1)} = \beta_k^{(t)} - \eta \left(2\lambda\beta_k^{(t)} + \sum_{i=1}^N (\mu_k(x_i) - 1\{y_i = k\}) x_i \right)$$

Stochastic gradient descent:

$$\beta_k^{(t+1)} = \beta_k^{(t)} - \eta \left(2\lambda\beta_k^{(t)} + (\mu_k(x_i) - 1\{y_i = k\}) x_i \right)$$

Question 2. Logistics Regression Consider the categorical learning problem consisting of a data set with two labels:

Label 1:

X_1	3.81	0.23	3.05	0.68	2.67
X_2	-0.55	3.37	3.53	1.84	2.74

Label 2:

X_1	-2.04	-0.72	-2.46	-3.51	-2.05
X_2	-1.25	-3.35	-1.31	0.13	-2.82

(1) Use **gradient descent** to find the logistic regression model

$$p(Y = 1|\vec{x}) = \frac{1}{1 + e^{-\beta^T \vec{x}}}$$

and the boundary. (Plot the boundary, only use numpy and Matplotlib.)

Using the following data matrix to do the logistic regression. (Notice that we need to use labels 0, 1)

X_1	X_2	Y
3.81	-0.55	1
0.23	3.37	1
3.05	3.53	1
0.68	1.84	1
2.67	2.74	1
-2.04	-1.25	0
-0.72	-3.35	0
-2.46	-1.31	0
-3.51	0.13	0
-2.05	-2.82	0

The (matrix notation) of the gradient of the Cross-Entropy cost J can be coded as

```
1 def sigmoid(x):
2     return 1/(1+np.exp(-x))
3
4 def grad_cost(theta, x, y):
5     z = x.dot(theta)
6     gradcost = (1/len(x))*np.matmul(x.T, (sigmoid(z)-y))
7     return gradcost
```

Define Gradient Descent function with iterations and learning rate alpha

```
1 def GradientDescent(x,y, theta, alpha, iteration):
2     for i in range(iteration):
3         theta_new = theta - alpha*grad_cost(theta,x,y)
4         theta = theta_new
5     return theta_new
```


The result of $\vec{\theta}$ depends on your initial value $\vec{\theta}_0$, number iterations, and learning rate α . With $\vec{\theta}_0 = \vec{0}$, $\alpha = 0.02$, and 1000 iterations, we get our $\vec{\theta}$:

[-0.04617983, -1.37920924, -1.25274956]


(Your answer may very different from this. But after divide θ_2 , the answer should be similarly. Or the boundary graph should be similarly.)

If you want, you can also recording the Cross-entropy cost values and plot them. The cross entropy function can be defined as:

```
1 def CELoss(x,y,theta):
2     z = y*x.dot(theta)
3     CE=np.matmul(y.T,np.log(sigmoid(z)))+np.matmul((np.ones(y.shape)-y).T,np.log((np.
4         ones(sigmoid(z).shape))))
5     return -(1/len(x))*CE
```



The boundary $\theta_0 + \theta_1 X_1 + \theta_2 X_2 = 0$ can be plotted using `plt.plot(X1, (-X1 * $\theta_1 - \theta_0$)/ θ_2 , color = "blue")` (Here, you only need to plot two points for X_1 , i.e, the min and the max.)



(2) Try **quadratic** Logistic Regression method for this question and obtain an quadratic boundary. (bonus)
 (Hint: this means to use new features: $X_1, X_2, X_1^2, X_1X_2, X_2^2$.)

Using the following data matrix to do the LDA again:

X_1	X_2	X_1^2	X_1X_2	X_2^2	Y
3.81	-0.55	3.81^2	$(3.81)(-0.55)$	$(-0.55)^2$	1
0.23	3.37				1
3.05	3.53	\vdots	\vdots	\vdots	1
0.68	1.84				1
2.67	2.74				1
-2.04	-1.25				0
-0.72	-3.35				0
-2.46	-1.31				0
-3.51	0.13				0
-2.05	-2.82				0

Redo the calculation based on the new data matrix. We have the $\vec{\theta}^T = [\theta_0 \theta_1 \theta_2 \theta_3 \theta_4 \theta_5]$:
 array([-0.01614066, -1.33955452, -1.23265001, 0.02176921, 0.20651087, -0.11120619])

(Again, your answer may very different from this. But after divide θ_2 , the answer should be similarly.
 Or the boundary graph should be similarly.)

The boundary $\theta_0 + \theta_1 X_1 + \theta_2 X_2 + \theta_3 X_1^2 + \theta_4 X_1 X_2 + \theta_5 X_2^2 = 0$


Remark: You may get the polynomial feature by basic coding: `numpy.c_[x, x1*x1, x1*x2, x2*x2]` to add columns. If allow to use scikit-learn in labs, we can use `sklearn.preprocessing` (See CVBootstrap.ipynb in lecture notes)

```
1 from sklearn.preprocessing import PolynomialFeatures
2 # Quadratic
3 poly = PolynomialFeatures(degree=2)
4 x_poly = poly.fit_transform(x)
5
```

Graphing: You may use the following code to draw the graph:

```
1 X, Y = np.meshgrid(np.arange(-4, 4, 0.05), np.arange(-4, 4, 0.05))
2 plt.contour(X, Y,
3 -0.01614066 - 1.33955452 * X - 1.23265001 * Y + 0.02176921 * X * X + 0.20651087 * X * Y - 0.11120619 * Y * Y,
4 [0])
5 plt.show()
```

The same drawing in different ranges.



Question 3. - Linear Discriminant Analysis: Consider the categorical learning problem consisting of a data set with two labels:

Label 1:

X_1	3.81	0.23	3.05	0.68	2.67
X_2	-0.55	3.37	3.53	1.84	2.74

Label 2:

X_1	-2.04	-0.72	-2.46	-3.51	-2.05
X_2	-1.25	-3.35	-1.31	0.13	-2.82

a) For each label above, the data follow a multivariate normal distribution $\text{Normal}(\mu_i, \Sigma)$ where the covariance Σ is the same for both label 1 and for label 2. Fit a pair of Gaussian discriminant functions to the labels by computing the covariances, means, and proportions of datapoints as discussed in the Linear Discriminant Analysis section. You may use a computer, but you should **not** use an LDA solver. You should report the values for μ_i and Σ .

$$\mu_1 = \begin{pmatrix} 2.088 \\ 2.186 \end{pmatrix}, \mu_2 = \begin{pmatrix} -2.156 \\ -1.72 \end{pmatrix}.$$

$$\Sigma = \frac{1}{10-2} \sum_{i=1}^{10} (X^{(i)} - \mu_k)(X^{(i)} - \mu_k)^T = \begin{pmatrix} 1.709575 & -1.23013 \\ -1.23013 & 2.349865 \end{pmatrix}.$$

$$\phi_1 = \phi_2 = \frac{5}{10} = \frac{1}{2}.$$

$$P(X|\text{Label} = 1) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1)\right)$$

$$P(X|\text{Label} = 2) = \frac{1}{(2\pi)^{n/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_2)^T \Sigma^{-1}(x - \mu_2)\right)$$

Suppose you already have the standard data matrices M1 and M2 as our standard form. (each one is a 5 by 2 matrix) The code for mean and covariance can be:

```
1 # means
2 mu1=M1.mean(0)
3 mu2=M2.mean(0)
```

```

4
5 # covariance matrix
6 (1/(len(M1)+len(M2)-2))*(np.matmul((M1-M1.mean(0)).T, M1-M1.mean(0))\
7 +np.matmul((M2-M2.mean(0)).T, M2-M2.mean(0)))
8
9 array([[ 1.709575, -1.23013 ],
10        [-1.23013 ,  2.349865]])

```

b) Give the **formula for the line** forming the discretion boundary.


The line forming the discretion boundary is $X = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$ such that $\log \frac{P(\text{Label}=2|X)}{P(\text{Label}=1|X)} = 0$.

$$P(\text{Label} = k|X) = \frac{P(X|\text{Label}=k)P(\text{Label}=k)}{P(X)}.$$

$$\begin{aligned} \log P(\text{Label} = k|X) &= \log P(X|\text{Label} = k) + \log P(\text{Label} = k) - \log P(X) \\ &= -\frac{1}{2} \log |\Sigma| - \frac{1}{2} (X - \mu_k)^T \Sigma^{-1} (X - \mu_k) + \log \phi_k + \text{constant} \\ &= X^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \phi_k - \frac{1}{2} \log |\Sigma| + \text{constant} \end{aligned}$$

Hence, $X = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$ such that $\log P(\text{Label} = 1|X) = \log P(\text{Label} = 2|X)$.

$$\begin{aligned} X^T \Sigma^{-1} \mu_1 - \frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 &= X^T \Sigma^{-1} \mu_2 - \frac{1}{2} \mu_2^T \Sigma^{-1} \mu_2 \\ (x_1 \ x_2) \begin{pmatrix} 3.03331817 \\ 2.51817686 \end{pmatrix} - 5.91915147956369 &= (x_1 \ x_2) \begin{pmatrix} -2.86820579 \\ -2.23343298 \end{pmatrix} - 5.012678200684864 \\ (x_1 \ x_2) \begin{pmatrix} 3.03331817 + 2.86820579 \\ 2.51817686 + 2.23343298 \end{pmatrix} &= -5.012678200684864 + 5.91915147956369 \\ 5.90152396x_1 + 4.75160984x_2 &= 0.9064732788788268 \\ x_2 &= \frac{0.9064732788788268 - 5.90152396x_1}{4.75160984} \end{aligned}$$



c) Use the **QDA** method for this question and obtain an quadratic boundary. (Hint, you need to calculate Σ_1 and Σ_2 separately.)

We assume the covariance Σ_1 and Σ_2 for each label are different. In this case,

$$\Sigma_1 = \frac{1}{5-1} \sum_{i=1}^5 (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T = \begin{bmatrix} 2.41602 & -1.202185 \\ -1.202185 & 2.78013 \end{bmatrix}$$

$$\Sigma_2 = \frac{1}{5-1} \sum_{i=1}^5 (x^{(i)} - \mu_{y^{(i)}})(x^{(i)} - \mu_{y^{(i)}})^T = \begin{bmatrix} 1.00313 & -1.258075 \\ -1.258075 & 1.9196 \end{bmatrix}$$


Simplify the calculation with constants, we have the equality

$$-\frac{1}{2} \log |\Sigma_1| - \frac{1}{2} (\vec{x} - \mu_1)^T \Sigma_1^{-1} (\vec{x} - \mu_1) + \log(\phi_1) = -\frac{1}{2} \log |\Sigma_2| - \frac{1}{2} (\vec{x} - \mu_2)^T \Sigma_2^{-1} (\vec{x} - \mu_2) + \log(\phi_2)$$

Plug in the information from (1) and Σ_1, Σ_2 we have the quadratic curve

$$2.536x^2 + 3.441xy + 19.982x + 1.234y^2 + 14.421y + 26.296 = 0$$

```
1 # covariance matrices
2 def covar(x):
3     return (1/(len(x)-1))*np.matmul((x-x.mean(0)).T, x-x.mean(0))
4
5 sigma_1 = covar(M1)
6 sigma_2 = covar(M2)
```



To simplify the formula, I used the **sympy** library. So I don't have to simplify by hand.

```
1 def quForm(x, S):
2     return np.matmul(np.matmul(x.T, S), x)
3
4 from sympy import *
5
6 xx, yy = symbols("xx yy")
7
8 simplify(math.log(phi_1)-math.log(phi_2)\
9 -0.5*quForm((np.array([xx,yy])-mu1), np.linalg.inv(sigma_1))\
10 +0.5*quForm((np.array([xx,yy])-mu2), np.linalg.inv(sigma_2))\
11 -0.5*math.log(np.linalg.det(sigma_1)) + 0.5*math.log(np.linalg.det(sigma_2))).evalf(4))
```

(d) Try quadratic LDA method for this question and obtain an quadratic boundary. (bonus)

Using the following data matrix to do the LDA again:

X_1	X_2	X_1^2	X_1X_2	X_2^2	Y
3.81	-0.55	3.81^2	$(3.81)(-0.55)$	$(-0.55)^2$	1
0.23	3.37				1
3.05	3.53	\vdots	\vdots	\vdots	1
0.68	1.84				1
2.67	2.74				1
-2.04	-1.25				0
-0.72	-3.35				0
-2.46	-1.31				0
-3.51	0.13				0
-2.05	-2.82				0

and obtain the formula

$$\vec{x}^T \Sigma^{-1} \vec{\mu}_1 - \frac{1}{2} \vec{\mu}_1^T \Sigma^{-1} \vec{x} + \log \phi_1 = \vec{x}^T \Sigma^{-1} \vec{\mu}_2 - \frac{1}{2} \vec{\mu}_2^T \Sigma^{-1} \vec{x} + \log \phi_2$$

Notice that our calculation is in dimension 5.

$$-2.34x^2 - 3.187xy + 20x - 0.914y^2 + 15y + 26.223 = 0$$

Again, we can use `plt.contour()` to draw the graph.

```

1 X,Y = np.meshgrid(np.arange(-50, 50, 0.05),np.arange(-50, 50, 0.05))
2 plt.contour(X,Y, -2.34*X*X-3.187*X*Y+20*X-0.914*Y*Y+15*Y+26.223, [0], colors="blue")
3
4 plt.legend(loc='lower right')
5 ax1.set_xlabel('X1')
6 ax1.set_ylabel('X2')
7 fig.set_size_inches(10, 7)
8 plt.show()

```

