

Lab 1 Exploratory Analysis - Problems (Sakshi Suman)

February 8, 2022

1 Exploratory Analysis

1.1 Problems:

Load the NYC AirBnB Truncated Dataset. This dataset is a mirror of the full NYC AirBnB dataset found at Kaggle, but only contains the first 10,000 entries.

<https://www.kaggle.com/dgomonov/new-york-city-airbnb-open-data>

For the numerical features,

- 1) Display histograms for the numerical features.
- 2) Construct the scatter plots of price with each of the numerical features.
- 3) Display the correlation histogram.
- 4) Using numerical features to predict the renting price.
- 5) Write down the predict function from (4)
- 6) Calculate the RSS cost.

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt

airbnb_data = pd.read_csv("https://raw.githubusercontent.com/tipthederiver/
↪Math-7243-2020/master/Datasets/NYCAirBnB/train.csv")
```

```
[2]: airbnb_data.head()
```

```
[2]:
```

	name	host_id	host_name	\
0	Clean & quiet apt home by the park	2787	John	
1	Skylit Midtown Castle	2845	Jennifer	
2	THE VILLAGE OF HARLEM...NEW YORK !	4632	Elisabeth	
3	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	
4	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	

	neighbourhood_group	neighbourhood	latitude	longitude	room_type	\
0	Brooklyn	Kensington	40.64749	-73.97237	Private room	

1	Manhattan	Midtown	40.75362	-73.98377	Entire home/apt
2	Manhattan	Harlem	40.80902	-73.94190	Private room
3	Brooklyn	Clinton Hill	40.68514	-73.95976	Entire home/apt
4	Manhattan	East Harlem	40.79851	-73.94399	Entire home/apt

	price	minimum_nights	number_of_reviews	last_review	reviews_per_month	\
0	149	1	9	10/19/2018	0.21	
1	225	1	45	5/21/2019	0.38	
2	150	3	0	NaN	NaN	
3	89	1	270	7/5/2019	4.64	
4	80	10	9	11/19/2018	0.10	

	calculated_host_listings_count	availability_365
0	6	365
1	2	355
2	1	365
3	1	194
4	1	0

```
[3]: airbnb_data.shape
```

```
[3]: (9999, 15)
```

```
[4]: airbnb_data.dtypes
```

```
[4]: name                object
     host_id             int64
     host_name           object
     neighbourhood_group object
     neighbourhood       object
     latitude            float64
     longitude           float64
     room_type           object
     price               int64
     minimum_nights      int64
     number_of_reviews    int64
     last_review         object
     reviews_per_month   float64
     calculated_host_listings_count int64
     availability_365     int64
     dtype: object
```

```
[5]: # Dropping non-numeric columns
     numeric_data = airbnb_data.drop(columns=['name',
                                             'host_id',
                                             'host_name',
                                             'neighbourhood_group'],
```

```
'neighbourhood',  
'room_type',  
'last_review']])
```

```
[6]: numeric_data.isnull().sum()
```

```
[6]: latitude          0  
longitude          0  
price              0  
minimum_nights     0  
number_of_reviews  0  
reviews_per_month  1322  
calculated_host_listings_count  0  
availability_365    0  
dtype: int64
```

```
[7]: numeric_data.dropna(inplace=True)
```

```
[8]: numeric_data.isnull().sum()
```

```
[8]: latitude          0  
longitude          0  
price              0  
minimum_nights     0  
number_of_reviews  0  
reviews_per_month  0  
calculated_host_listings_count  0  
availability_365    0  
dtype: int64
```

```
[9]: numeric_data.dtypes
```

```
[9]: latitude          float64  
longitude          float64  
price              int64  
minimum_nights     int64  
number_of_reviews  int64  
reviews_per_month  float64  
calculated_host_listings_count  int64  
availability_365    int64  
dtype: object
```

```
[10]: numeric_data.shape
```


```
[10]: (8677, 8)
```

```
[11]: X_df = numeric_data.drop(columns='price')  
X = X_df.values
```

```
Y_df = numeric_data['price']
Y = Y_df.values
X = np.concatenate((np.ones(X.shape[0]).reshape((-1, 1)), X), axis=1)
```

2 Histogram

```
[12]: numeric_data.hist(bins=50, figsize=(40, 30))
plt.show()
```



```
[13]: column_names = list(X_df.columns)
column_names
```


```
[13]: ['latitude',
       'longitude',
       'minimum_nights',
       'number_of_reviews',
       'reviews_per_month',
       'calculated_host_listings_count',
       'availability_365']
```

3 ScatterPlot

```
[14]: f, axes = plt.subplots(1, 7, sharey=True)
f.set_size_inches(14, 2)
f.tight_layout()

axes = axes.reshape(7)


for i in range(len(column_names)):
    axes[i].plot(X[:, i + 1], Y, 'x', color='Blue')
    axes[i].set_title(column_names[i], fontsize=12)
```



4 Correlation

```
[15]: correlation_matrix = X_df.corr()
```

```
[16]: fig, ax = plt.subplots(figsize=(14,10))
sns.heatmap(correlation_matrix, ax=ax, linewidths=0.05, cmap="magma", annot=True)
plt.show()
```



5 Linear Regression

```
[17]: # Linear regression matrix calculation
def normal_equation(x, y, w=None):
    if w is None:
        return np.linalg.inv(x.T.dot(x)).dot(x.T).dot(y)
    else:
        return np.linalg.inv(x.T.dot(w).dot(x)).dot(x.T).dot(w).dot(y)
```

```
[18]: theta = normal_equation(X, Y)
theta
```

```
[18]: array([-6.82277031e+04,  1.28955835e+02, -8.53509860e+02,  2.45529700e-02,
          -1.34321164e-01, -4.64585556e+00, -8.99817604e-01,  1.32849443e-01])
```

6 Predict function

```
[19]: def predict(x, theta):  
       return np.dot(x, theta)
```

```
[20]: predict(X[0, :], theta)
```

```
[20]: 191.10624136681963
```

7 RSS

```
[21]: def rss(X, Y):  
       theta = normal_equation(X, Y)  
       error = X @ theta - Y  
       return np.linalg.norm(error) / Y.shape[0]
```

```
[22]: rss(X, Y)
```

```
[22]: 2.569281936431202
```